

Samrat Sahoo

972-971-9060 | samrat@stanford.edu | [linkedin.com/in/samratsahoo/](https://www.linkedin.com/in/samratsahoo/) | github.com/SamratSahoo

EDUCATION

Stanford University (Incoming)

Master of Science in Computer Science

September 2025 – June 2027

Stanford, CA

Georgia Institute of Technology

Bachelor of Science in Computer Science, Minor in Economics | GPA: 3.97/4.00

Aug. 2021 – May 2025

Atlanta, GA

- Concentrations: Artificial Intelligence + Systems and Architecture
- Awards: Bessemer Fellow, FinTech Fellow, Faculty Honors, Startup Pitch Competition Winner
- Coursework: Data Structures & Algorithms, Operating Systems, Compilers and Interpreters, Deep Learning

RESEARCH PROJECTS & PUBLICATIONS

Financial Portfolio Optimization Using Deep Reinforcement Learning | [Project Website](#)

- Created an empirically-driven reinforcement learning environment to optimize retirement portfolios, simulating real-world circumstances better than synthetic data-driven environments
- Trained a soft actor-critic agent to create an optimal portfolio based on a user-defined risk-reward threshold

Exploration Critique Networks | [Project Website](#)

- Leveraging cross-attention between past trajectories and recent state transitions to guide RL exploration
- Implemented state compression using autoencoders to learn useful state representations by predicting current and future states with an augmented loss function using episodic reward

Generative Offline Exploration using Decision Transformers | [Project Website](#)

- Developed a novel data augmentation approach for offline reinforcement learning using Decision Transformers to generate synthetic trajectories in data-scarce environments.
- Trained a behavioral cloning and DDPG agent with and without synthetic trajectories to measure performance

Scatter Protocol | *Accepted and published at IEEE International Conference on Blockchain (Blockchain 2024)*

- Developed an incentivized and trustless protocol for decentralized, federated learning (Solidity/Go-Ethereum)
- Created a decentralized, peer2peer node system to split machine learning loads across a network (GoLang/LibP2P)
- Designed a secure machine learning environment by introducing OS isolation using Open Container Initiative
- **Citation:** S. Sahoo and S. Chava, "Scatter Protocol: An Incentivized and Trustless Protocol for Decentralized Federated Learning," 2024 IEEE International Conference on Blockchain (Blockchain), Copenhagen, Denmark, 2024, pp. 497-504, doi: 10.1109/Blockchain62396.2024.00073.

Teaching ChatGPT | *Accepted and published at Human Factors and Ergonomics Society*

- Directed experiments with 150+ students to determine the efficacy of the protege effect applied with LLMs
- Conducted a literature review regarding teachable agents and psychological concomitants of the protege effect
- **Citation:** Tidler, Z. R., Sahoo, S., & Catrambone, R. Teaching ChatGPT: Attempting to Demonstrate the Protégé Effect with a Large Language Model Learner. In *Proceedings of the 68th International Annual Meeting of the Human Factors and Ergonomics Society*. Phoenix.

Universal NFT Vector Database | *Preprint*

- Developed public visualization and similarity matching APIs (Flask/Google Cloud Run) using RegNetY-080
- Devised a public graphical user interface (Vue.js/MongoDB) for easy interaction with vector database data
- Integrated task queue structure (AWS SQS/Celery) for horizontal scaling and hardware agnosticism
- **Citation:** S. Sahoo, N. Paul, A. Shah, A. Hornback, S. Chava. "The Universal NFT Vector Database: A Scalable Vector Database for NFT Similarity Matching." (2023).

RESEARCH EXPERIENCE

Research Assistant

Georgia Tech Financial Services and Innovation Lab

December 2021 – May 2025

Atlanta, GA

- Designed and developed a large-scale database (Vue.js/Celery/Python) for Digital Asset counterfeit detection
- Built a privacy-centered federated learning protocol for collaborative machine learning (GoLang/Python/LibP2P)
- Leveraging deep reinforcement learning techniques to optimize retirement portfolio investments (PyTorch/NumPy)

WORK EXPERIENCE

Software Engineering Intern

May 2023 – July 2023

Cruise Automation (YC W14)

San Francisco, CA

- Refactored the access control system, saving 100+ hours of engineering time per month (PostgreSQL/OPA)
- Engineered widgets displaying exit codes, and execution cost, saving 1000s of dollars (Kubernetes/Pulumi)
- Extended simulation platform's search engine to drive additional insights into tests (ElasticSearch/BigQuery)

Software Engineering Intern

Jan. 2023 – May 2023

Visor

New York City, NY

- Redeveloped the HubSpot chat integration for better user experiences for 1000+ users (Vue.js/HubSpot API)
- Built the settings integrations page from scratch for different CRMs (HubSpot/Salesforce/Atlassian APIs)
- Redesigned the Visor home page, optimizing for real-time filtering via graph manipulation (CloudstoreDB)

Software Engineering Intern

June 2022 – August 2022

Fidelity Investments

Westlake, TX

- Implemented a data interaction application (Vue.js/Express.js/AWS), saving 100+ hours of manual data mining
- Delivered automation APIs (GraphQL/Python) to 100 teams to interact with legacy engineering infrastructure
- Deployed data mining applications to 3000+ developers (Jenkins/ uDeploy), improving integration testing time
- Built modern GraphQL API layers on top of legacy SOAP APIs, reducing development time by 66 percent

Software Engineering Intern

March 2021 – August 2021

Roboflow (YC S20)

Des Moines, IA

- Created the training procedure for the Roboflow classification network (PyTorch/Docker) for 3 enterprise users
- Initiated the Roboflow SDK (Python) enabling better workflow integrations for 350,000+ developers
- Roboflow Python package integrated with YOLOv5 (25,000 stars on Github), increasing users by 1000 per month
- Adapted vision models like YOLOX and MobileNetV2 (C++/Python) for Roboflow (1000+ monthly users)

OPEN SOURCE SOFTWARE ENGINEERING PROJECTS

Bits of Good Unified Analytics Platform

- Devised an analytics dashboard for viewing analytics and managing projects for all 150 members at Bits of Good
- Deployed the API using serverless infrastructure to scale to economically handle 500,000 events per month
- Leveraged cron jobs through GitHub Actions to periodically ping analytics services for health monitoring purposes

Roboflow Python Package

- Extended the API to add better support to package functionality and edge cases for better developer experiences
- Architected the design of the Python package, establishing the direction of future extension and development
- Published to the package directory, resulting in 500,000+ monthly downloads in top machine learning repositories

LEADERSHIP

Director of Engineering, Engineering Manager

September 2021 – Present

Bits of Good

Atlanta, GA

- Led a team of developers using agile methodology to create a mobile app (React Native) for over 800 users
- Connected and streamlined 1400 developers through creating the National Hack4Impact Portal (Next.js/Azure)
- Began the Infrastructure Labs team to build scalable systems to reduce total development time by 100s of hours

TECHNICAL SKILLS

- **Languages and Frameworks:** PyTorch, NumPy, Python, Javascript, Java, Solidity, Typescript, GoLang, React.js, Vue.js, Nuxt.js, Next.js, Express.js, Flask, React.js, MongoDB, React Native, Assembly, C, LibP2P, SQL, Gymnasium (Reinforcement Learning)
- **Developer Tools and Libraries:** Git, Amazon Web Services (Simple Queue Service, Elastic Container Service), Google Cloud Product (Cloud Run), REST API, Docker, Docker Compose, GitHub Actions, Linux, CLI, RabbitMQ, Firebase, OpenAI Vector Store, Postman